

Chapter 3: Processes





Process Concept


- A **process** is a program in execution
- program vs process
 - A program is a **passive** entity stored on disk (**executable file**)
 - A process is an **active** entity
 - A program becomes a process when executable file is loaded into memory





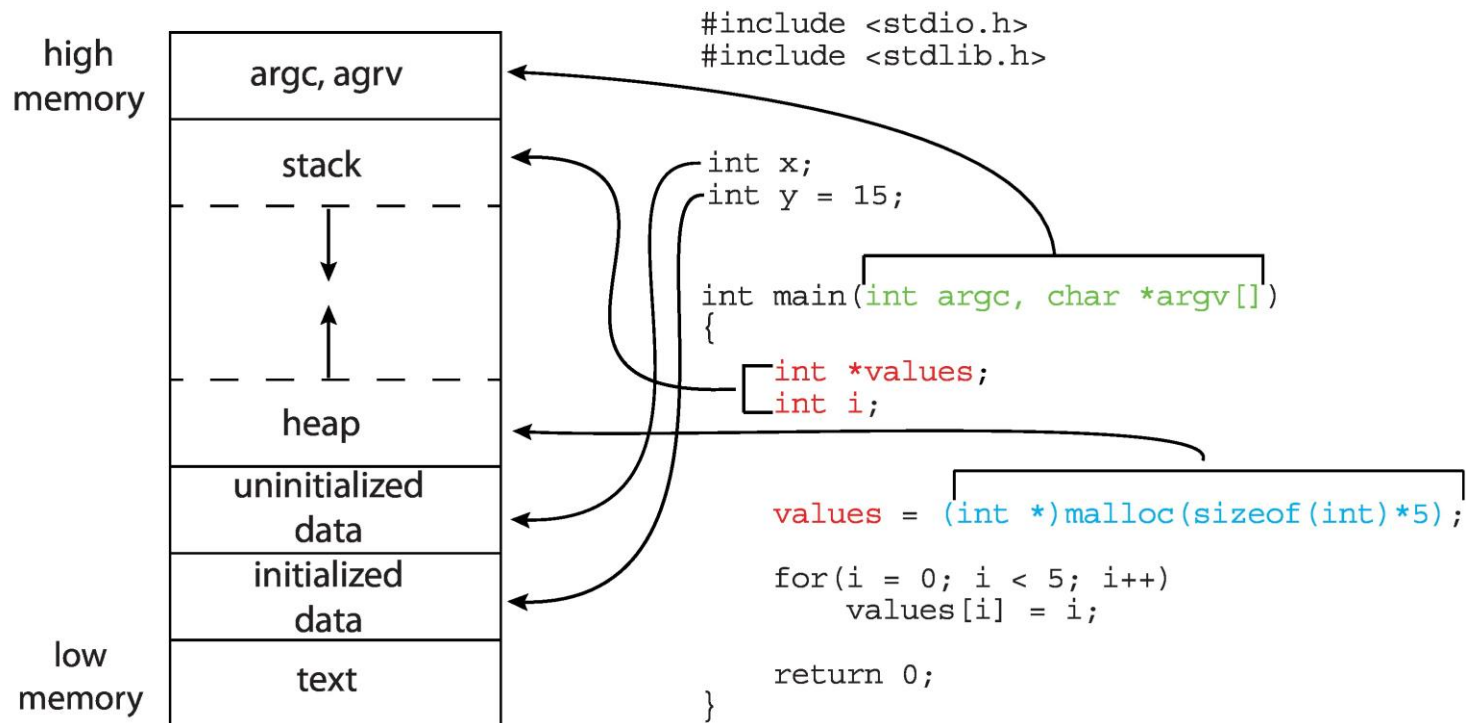
Parts of a Process

- The address space of a process consists of
 - **Text section** containing the program code
 - **Data section** containing global variables
 - **Stack** containing temporary data when invoking functions
 - ▶ parameters, return addresses, local variables
 - **Heap** containing memory dynamically allocated during run time
- Current activity of a process represented by
 - Value of **program counter** (i.e., address of next instruction to execute)
 - Content of CPU registers





Memory Layout of a C Program



[Python Tutor link](#)
Size a.out





Relationship between Processes and Programs

- Multiple processes may be associated with the same program
 - The processes are separate execution sequences
 - The processes have the same text section
 - The data, heap, and stack sections vary
 - Value of program counter and contents of CPU registers also vary





Process State

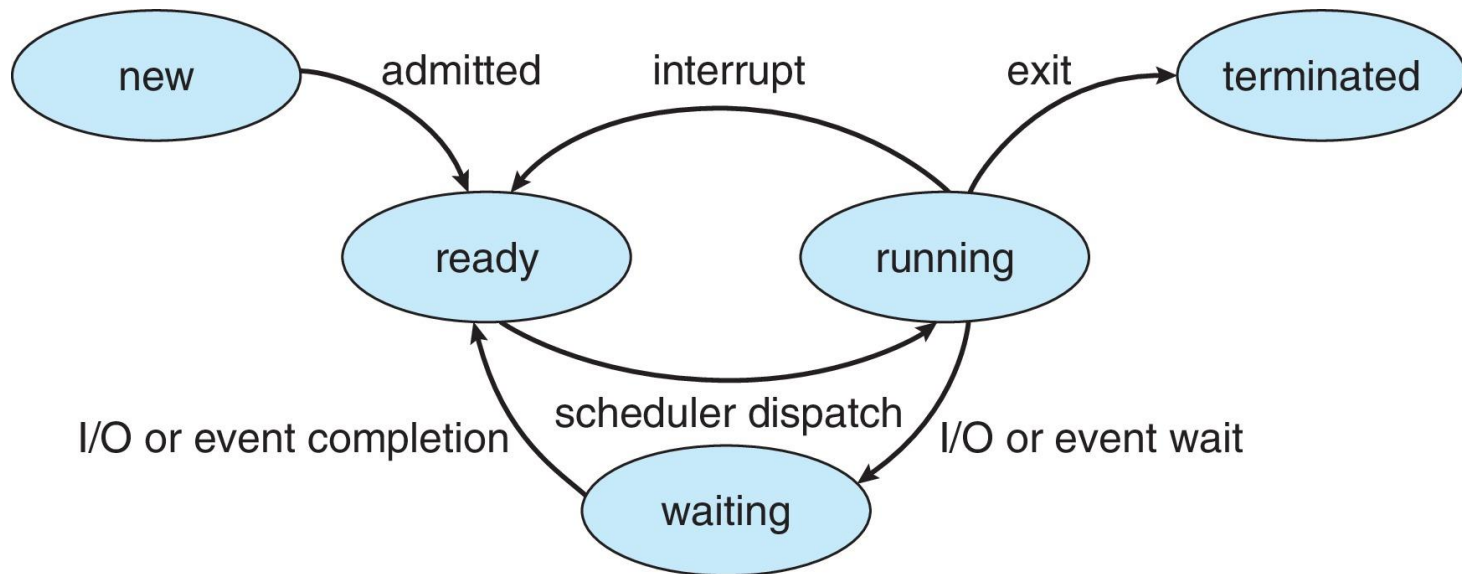
- As a process executes, it changes **state**
 - **New**: The process is being created
 - **Running**: Instructions are being executed
 - **Waiting**: The process is waiting for some event to occur (e.g., I/O completion)
 - **Ready**: The process is waiting to be assigned to a processor
 - **Terminated**: The process has finished execution
- Only one process can be **running** on a processor at any instant
- Many processes may be **ready** and **waiting**





State Transitions

A **state transition** for a process is a change in its state caused by the occurrence of some event



Show Animation 3.1



Process Control Block (PCB)

Process control block (PCB) stores information associated with each process

- ❑ **Process state** – running, waiting, etc
- ❑ **Process number** – unique ID assigned to process
- ❑ **Program counter** – address of next instruction to execute
- ❑ **CPU registers** – contents of all CPU registers
- ❑ **CPU scheduling information**- priorities, scheduling queue pointers
- ❑ **Memory-management information** – memory allocated to the process
- ❑ **Accounting information** – CPU time used, clock time elapsed since start, time limits
- ❑ **I/O status information** – I/O devices allocated to process, list of open files

process state
process number
program counter
registers
memory limits
list of open files
...





Process Scheduling

- Operating systems use multiprocessing to maximize CPU utilization
- Operating systems use multitasking to enable users to interact with their programs
 - OS switches a CPU core among processes frequently
- **Process scheduler** selects among **ready** processes for next execution on CPU core
- OS maintains **scheduling queues** of processes
 - **Ready queue** – set of processes residing in main memory, ready and waiting to execute
 - **Wait queues** – set of processes waiting for an event (i.e. I/O completion)






Process Representation in Linux

- [linux/sched.h](#)
- Struct task_struct*
- Doubly linked list



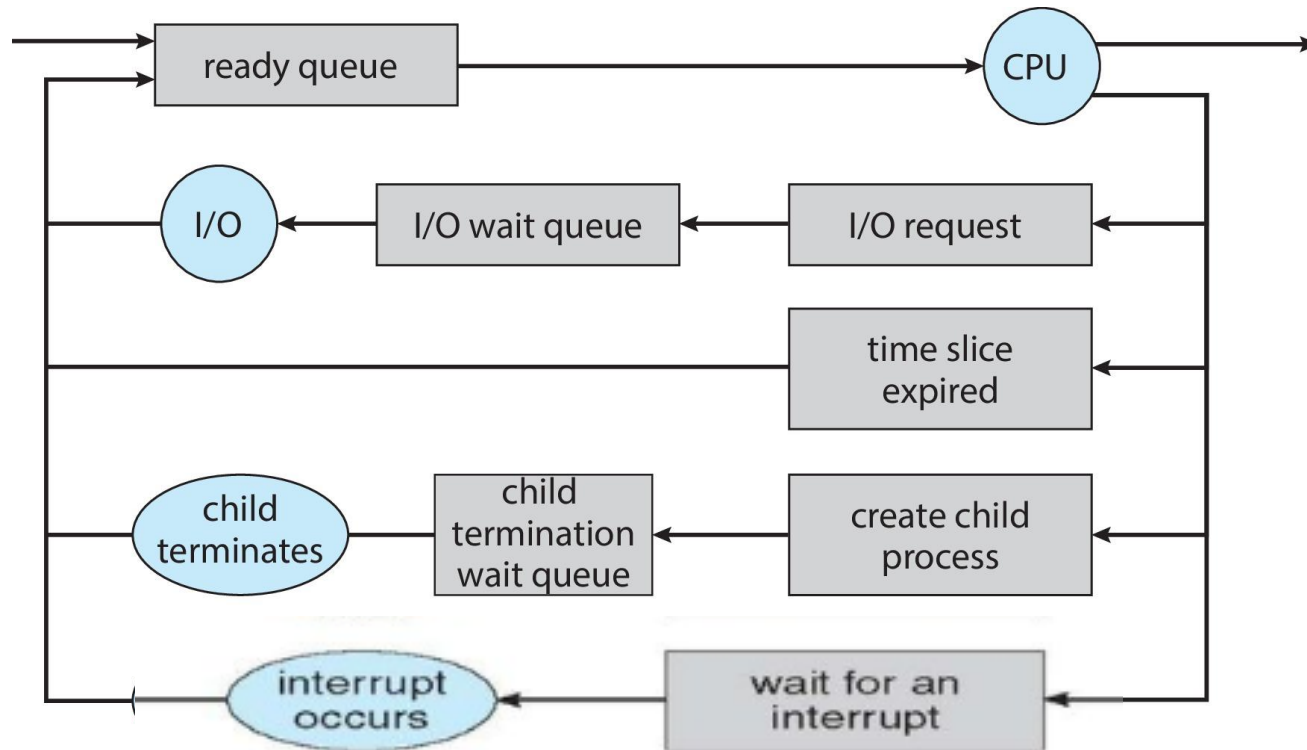


Ready and Wait Queues





Representation of Process Scheduling



Processes migrate among the queues





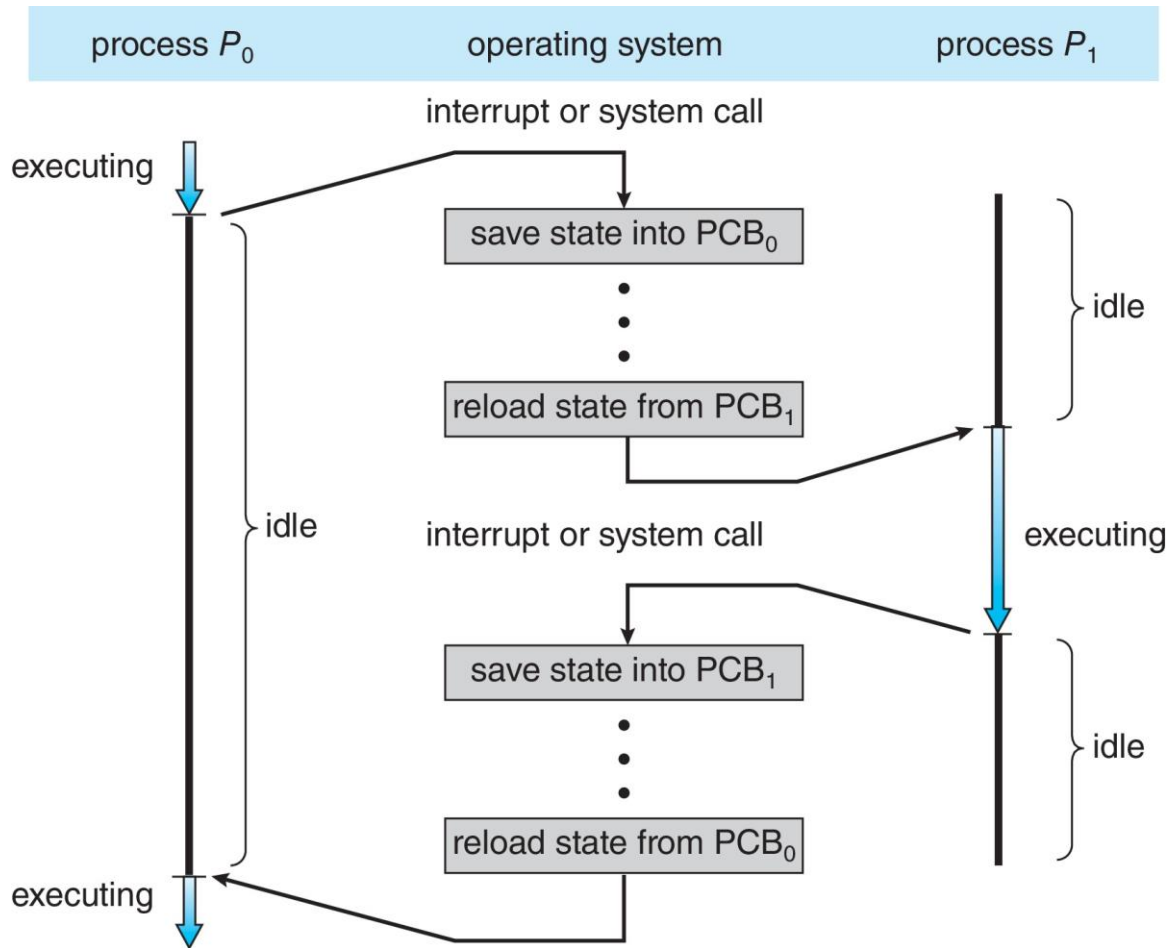
Context Switch

- ❑ When CPU switches to another process, the system must **save the context** of the old process and load the **saved context** for the new process via a **context switch**
- ❑ **Context** of a process saved in its PCB
 - ❑ Value of CPU registers, program counter, memory management-information
- ❑ Context-switch time (typically a few microseconds) is **overhead**; the system does no useful work while switching





CPU Switch From Process to Process



Show Animation 3.2





CPU Efficiency in Multitasking Systems

- In a multitasking system, CPU switches rapidly between processes
- A **time slice**, denoted by δ , is the largest amount of CPU time any process can consume when scheduled to execute on the CPU
- If time slice elapses before the process completes its execution, kernel preempts the process
- Let σ denote the context-switch time, then CPU efficiency = $\delta / (\delta + \sigma)$





Operations on Processes

- OS must provide mechanisms for:
 - process creation
 - process termination





Process Tree in Linux

- ❑ Show `ps -el`
- ❑ Process tree all the way upto system (eg: `ps`)
- ❑ <https://en.wikipedia.org/wiki/Systemd>
- ❑ Figure 3.7: Linux Process Tree





Process Creation

- ❑ **Parent** process creates **child** processes, which, in turn create other processes, forming a **tree** of processes
 - ❑ Each process identified by a **process identifier (pid)**
- ❑ Resource sharing options
 - ❑ Parent and children share all resources
 - ❑ Child shares subset of parent's resources
 - ❑ Parent and child share no resources
- ❑ Address-space possibilities for the child process:
 - ❑ Child is a duplicate of parent (child has same program and data as parent)
 - ❑ Child has a new program loaded into it
- ❑ Execution options
 - ❑ Parent and child execute concurrently
 - ❑ Parent waits until child terminates





fork() system call

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t fork(void); // create a child process identical to parent
```

- Child process has a copy of parent's address space
- On success:
 - Both parent and child continue execution at the instruction after fork()
- On failure: Returns pid of the child process to the parent process; returns 0 to child process
 - Child is not created; returns -1 to parent





fork() example

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>

int main()
{
    int x=5;
    pid_t pid = fork();
    if (pid < 0) return -1;
    printf("I am running\n");
    if (pid) {
        printf("I am the parent\n");
        x *= 2;
    } else {
        printf("I am the child\n");
        x += 2; }
    printf("x is %d\n", x);
    return 0;
}
```

Output

```
P: I am running
P: I am the parent
P: x is 10
C: I am running
C: I am the child
C: x is 7
```





Fork() Example 2

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
int main()
{
    int x=5;
    pid_t pid = fork();
    if (pid < 0) return -1;
    printf("I am running\n");
    if (pid) {
        x = 10;
        pid_t pid = fork();
        printf("I am the parent\n");
        x *= 2;
    } else {
        printf("I am the child\n");
        x += 2; }
    printf("x is %d\n", x);
    return 0;}
```

P0: I am running
P1: I am running
P1: I am the child
P1: x is 7
P2: I am the parent
P2: x is 20
P0: I am the parent
P0: x is 20





Process Termination (1)

- Process executes last statement and then asks the OS to delete it using the `exit()` function.
 - Process' resources (e.g., memory, open files) are deallocated by OS

`#include <stdlib.h>`

`void exit(int status);` //cause normal process termination

- The function does not return
- Parent calls `wait()` to obtain status value from child





Process Termination (2)

- Parent process may wait for termination of a child process by using the `wait()` system call
- `wait()` returns status information and the pid of the terminated process

```
pid = wait(&status); //status is an integer pointer
```
- When a process terminates, its PCB must remain until parent calls `wait()` as process' exit status is stored in its PCB
- If a process has terminated, but parent has not yet called `wait()`, then process becomes a **zombie**
- If parent terminated without invoking `wait()`, then the process becomes an **orphan**
 - UNIX assigns the `init` process as the new parent to orphan processes; `init` process periodically invokes `wait()`
 - In Linux, `systemd` process becomes the parent to orphan processes





Waiting for Process Termination (1)

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

pid_t wait(int *status) – wait for any child to terminate

- If **status** is not NULL, the exit status of child is stored in the location pointed to by **status**
- On success, returns the pid of the child that terminated; on error, -1 is returned





Waiting for Process Termination (2)

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

pid_t waitpid(pid_t pid, int *status, int options) – wait for the specified child to terminate

- Options
 - 0: wait until child has exited
 - WNOHANG: return immediately if child has not exited
- On success, returns the pid of the terminated child or zero if WNOHANG is used and the specified child has not terminated
- On error, -1 is returned

Note: **wait(&status)** is equivalent to **waitpid(-1, &status, 0)**





wait() example

```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main()
{
    pid_t pid = fork();
    if (pid < 0) return -1;
    if (pid==0) {
        printf("I am the child\n");
        exit(0);
    }
    waitpid(pid, NULL, 0);
    printf("Child has terminated, so will I\n");
    return 0;
}
```

Output

C: I am the child

P: Child has terminated, so will I





Executing Another Program from a Process (1)

```
#include <unistd.h>
```

```
int execlp(const char *file, const char *arg, ...)
```

```
int execvp(const char *file, char *const argv[ ])
```


- The functions are used to replace the current process image with a new process image.
- The first argument is the name of a file that is to be executed
- **const char *arg, ...** in `execlp()` describes a list of one or more pointers to null-terminated strings that represent the argument list to the executed program
- **char *const argv[]** in `execvp()` is an array of pointers to null-terminated strings that represent the argument list to the executed program
- The first argument in argument list/array should point to the filename associated with the file to be executed
- The argument list/array must be terminated by a NULL pointer
- The functions do not return on success and return -1 on error





Executing Another Program from a Process (2)

- ❑ Parent uses **fork()** to create a child process
- ❑ Child uses **exec()** to replace its memory space with a new program
- ❑ Parent calls **wait()** to wait for child to terminate





execvp() Example

```
/*
This program will
1) Create a process
2) Have the new process execute "ls -a -l"
   using execvp()
3) Wait until the new process terminates
*/

#include <sys/types.h>
#include <sys/wait.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        printf("Fork Failed\n");
        exit(1);
    }
    else if (pid == 0) { /* child process */
        execvp("ls", "ls", "-a", "-l", NULL);
        // still here? Problem...
        printf("Exec error!\n");
        exit(1);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        waitpid(pid, NULL, 0);
        printf("Child has terminated, so will I\n");
    }
    exit(0);
}
```





execvp() Example

```
/*  
    This program will  
  
    1) Create a process  
    2) Have the new process execute "ls -a -l"  
        using execvp()  
    3) Wait until the new process terminates  
  
*/
```

```
#include <sys/types.h>  
#include <sys/wait.h>  
#include <stdio.h>  
#include <stdlib.h>  
#include <unistd.h>
```

```
int main()  
{  
    pid_t pid;  
  
    /* fork a child process */  
    pid = fork();  
    if (pid < 0) { /* error occurred */  
        printf("Fork Failed\n");  
        exit(1);  
    }  
    else if (pid == 0) { /* child process */  
        char *args[4]={"ls", "-a", "-l", NULL};  
        execvp("ls", args);  
        // still here? Problem...  
        printf("Exec error!\n");  
        exit(1);  
    }  
    else { /* parent process */  
        /* parent will wait for the child to complete */  
        waitpid(pid, NULL, 0);  
        printf("Child has terminated, so will \n");  
    }  
    exit(0);  
}
```





File System Calls

- `open()`: open a file
- `close()`: close a file
- `read()`: read data from a file
- `write()`: write data to a file





Open()

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
```

```
int open(const char *pathname, int flags, mode_t mode);
```

```
//open and possibly create a file
```

- Return value

- On success, returns a file descriptor (a small, nonnegative integer) for use in subsequent system calls
- On error, returns -1





An Open() Example

```
int fd = open("file.txt", O_WRONLY|O_CREAT|O_TRUNC,  
S_IRWXU|S_IROTH)
```

- **flags**

- O_WRONLY: write-only
- O_RDONLY: read-only
- O_RDWR: read/write
- O_CREAT: create the file if it does not exist
- O_TRUNC: if file already exists and access mode allows writing (i.e., O_RDWR or O_WRONLY), it will be truncated to length 0

- **mode** specifies the permissions to use in case a new file is created

- S_IRWXU: user (file owner) has read, write, and execute permission
- S_IROTH: others have read permission
- **mode** must be specified when O_CREAT is specified in flags, and is ignored otherwise.





Read(), Write(), Close()

#include <unistd.h>

ssize_t read(int fd, void *buf, size_t count) // read from a file descriptor

- Reads up to count bytes from file descriptor fd into the buffer pointed to by buf
- Returns the number of bytes read on success (zero indicates end of file) and -1 on error

ssize_t write(int fd, const void *buf, size_t count) // write to a file descriptor

- Writes up to count bytes from the buffer pointed to by buf to file descriptor fd
- Returns the number of bytes written on success and -1 on error

int close(int fd) //close a file descriptor


- Closes the file descriptor fd so that it no longer refers to any file and may be reused
- Returns 0 on success and -1 on error.





Pipes

- A pipe is a **unidirectional** data channel that can be used for interprocess communication
 - Sender puts data to one end (the **write-end** of the pipe)
 - Receiver gets data from the other end (the **read-end** of the pipe)





Creating Pipes

```
#include <unistd.h>
```

```
int pipe(int pipefd[2]) // create a pipe
```


- The array `pipefd` is used to return two file descriptors referring to the two ends of the pipe
 - `fd[0]` refers to the **read-end** of the pipe
 - `fd[1]` refers to the **write-end** of the pipe
 - Data written to the write end of the pipe is buffered by the kernel until it is read from the read end of the pipe.
- Returns 0 on success and -1 on error





Using Pipes (1)


- Sending process and receiving process must have a common ancestor who creates a pipe
- A child process inherits open files from parent, and a pipe is a special type of file → Child inherits pipe from parent





Using Pipes (2)

- Sending process calls `write(fd[1], buf, len)` to write to the pipe
- Receiving process calls `read(fd[0], buf, len)` to read from the pipe
- Each process should close the unused end of pipe to ensure that a process reading from the pipe can detect end-of-file when the writer has closed its end of the pipe
 - Sending process should close `fd[0]`
 - Receiving process should close `fd[1]`



Parent sends message to child

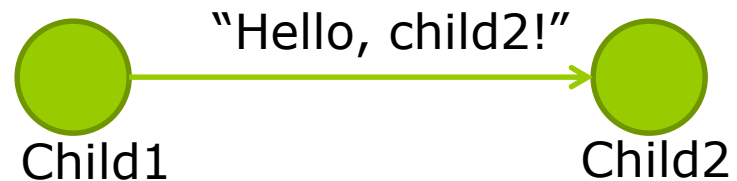




Pipe Example

```
/*
 * This program creates two child processes
 * and a pipe. Child 1 sends a message to child2
 * via the pipe.
 */
#include <sys/wait.h>
#include <unistd.h>
#include <string.h>
#include <stdio.h>
void Child1(int fd)
// fd is the file descriptor to write the message to
{
char *message = "Hello, child2!";
int bytes = strlen(message)+1;
printf("Child1 sending %d byte message %s to\n", bytes, message);
write(fd, message, bytes);
printf("Child1 exiting\n");
}
```

```
void Child2(int fd)
// fd is the file descriptor to read from
{
char buffer[256];
printf("Child2 waiting for message from\n", child1);
int sent = read(fd, buffer, 256);
printf("Child2 got %d byte message %s from\n", sent, buffer);
printf("Child2 exiting\n");
}
```

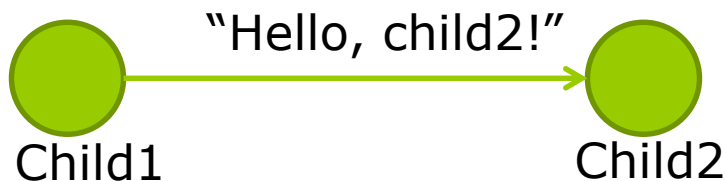




Pipe Example (continued)

```
int main()
{
    int pipefd[2];
    if (pipe(pipefd) < 0) {
        printf("Pipe error!\n");
        return 1;
    }
    pid_t pid2 = fork();
    if (pid2==0) {
        // Child 2
        close(pipefd[1]); //close write-end of pipe
        Child2(pipefd[0]); //read from pipe
        return 0;
    }
```

```
pid_t pid1 = fork();
if (pid1==0) {
    // Child 1
    close(pipefd[0]); //close read-end of pipe
    Child1(pipefd[1]); // write to pipe
    return 0;
}
// Parent... close pipe and wait for children
close(pipefd[0]);
close(pipefd[1]);
waitpid(pid1, NULL, 0);
waitpid(pid2, NULL, 0);
printf("Everyone is done\n");
return 0;
}
```





Redirection of Standard Input and Output

- In the UNIX shell, the symbols '<' and '>' cause input and output to be redirected
- For example
 - **ls -l > file1** sends the output of **ls -l** to file1
 - **head -5 < file1** reads from file1 instead of standard input
- We can use **dup2()** system call to redirect standard input and output





dup2()

```
#include <unistd.h>
```

```
int dup2(int oldfd, int newfd) //duplicate a file descriptor
```

- ❑ Makes newfd be a copy of oldfd
- ❑ Returns the new file descriptor, or -1 if an error occurred

- ❑ **Standard UNIX file descriptors:**
 - 0 = standard input (stdin)
 - 1 = standard output (stdout)
- ❑ **To redirect standard input:**

```
int fd = open ("input.txt", O_RDONLY);  
dup2(fd, 0); //whenever your program reads from stdin, it will read  
                from fd
```
- ❑ **To redirect standard output:**

```
int fd = open ("output.txt", O_WRONLY|O_TRUNC);  
dup2(fd, 1); //whenever your program writes to stdout, it will write to  
                fd
```





Pipeline Example

/* * This program:

- * Creates two children processes and a pipe.
 - * Re-directs standard output of child1 to the pipe.
 - * Re-directs standard input of child2 to the pipe.
 - * Runs "yes no" in child1
 - * Runs "head -10" in child2
 - * Waits for all to finish
- */

```
#include <stdio.h>
#include <sys/wait.h>
#include <unistd.h>
#include <string.h>
void Child1(int fd)
// fd is the file descriptor to write to
{
    dup2(fd, 1);
    close(fd);
    execlp("yes", "yes", "no", NULL);
    printf("Child 1 error!\n");
}
void Child2(int fd)
// fd is the file descriptor to read from
{
    dup2(fd, 0);
    close(fd);
    execlp("head", "head", "-10", NULL);
    printf("Child 2 error!\n");
}
```

```
int main()
{
    int pipefd[2];
    if (pipe(pipefd) < 0) {
        printf("Pipe error!\n");
        return 1;
    }
    int pid2 = fork();
    if (0==pid2) { // Child 2
        close(pipefd[1]); // I don't write to the pipe
        Child2(pipefd[0]); // read end of pipe
        return 0;
    }
    int pid1 = fork();
    if (0==pid1) { // Child 1
        close(pipefd[0]); // I don't read from the pipe
        Child1(pipefd[1]); // write end of pipe
        return 0;
    }
    // Parent... close pipe and wait for children
    close(pipefd[0]); // if you take this out, we will hang!
    close(pipefd[1]);
    waitpid(pid1, NULL, 0);
    waitpid(pid2, NULL, 0);
    printf("Everyone is done\n");
    return 0;
}
```





- We have gone through all the system calls you need to complete Project 1-UNIX Shell
- Start early!

