



# Announcements


---

- Midterm exam will be on 25 Sep Friday, 6:45pm (night exam)
- Study guide posted
- Exam Review 23-Sep
- Recitation will discuss homework solutions
- Project1 due Sep 20<sup>th</sup>, list of commands posted



# Chapter 6: Synchronization Tools

Section 6.1-6.7





# Background

---

- ❑ Processes can execute concurrently
  - ❑ May be interrupted at any time, partially completing execution before another process is scheduled
- ❑ Concurrent access to shared data may result in data inconsistency
- ❑ Maintaining data consistency requires mechanisms to ensure the orderly execution of **cooperating processes**
  - ❑ A **cooperating process** shares data with other processes, can affect or be affected by other processes
- ❑ Illustration of the problem:
  - ❑ A producer process & a consumer process share a buffer of certain size
  - ❑ A shared integer variable **counter** is used to keep track of the number of items in buffer
  - ❑ **counter** incremented by producer after it produces an item, decremented by consumer after it consumes an item





# Producer-Consumer Problem

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
}item;
item buffer[BUFFER_SIZE]; //shared variable
int counter = 0; //shared variable
```


```
int in = 0;
while (true) {
    /* produce an item in
       next_produced */
    while (counter == BUFFER_SIZE);
    /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

Producer code

```
int out = 0;
while (true) {
    while (counter == 0);
    /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    counter--;
    /* consume the item in
       next_consumed */
}
```

consumer code





# Race Condition

- **counter++** could be implemented in machine language as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter--** could be implemented in machine language as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with “counter = 5” initially:

S0: producer execute	<b>register1 = counter</b>	{register1 = 5}
S1: producer execute	<b>register1 = register1 + 1</b>	{register1 = 6}
S2: consumer execute	<b>register2 = counter</b>	{register2 = 5}
S3: consumer execute	<b>register2 = register2 - 1</b>	{register2 = 4}
S4: producer execute	<b>counter = register1</b>	{counter = 6}
S5: consumer execute	<b>counter = register2</b>	{ <b>counter = 4</b> }

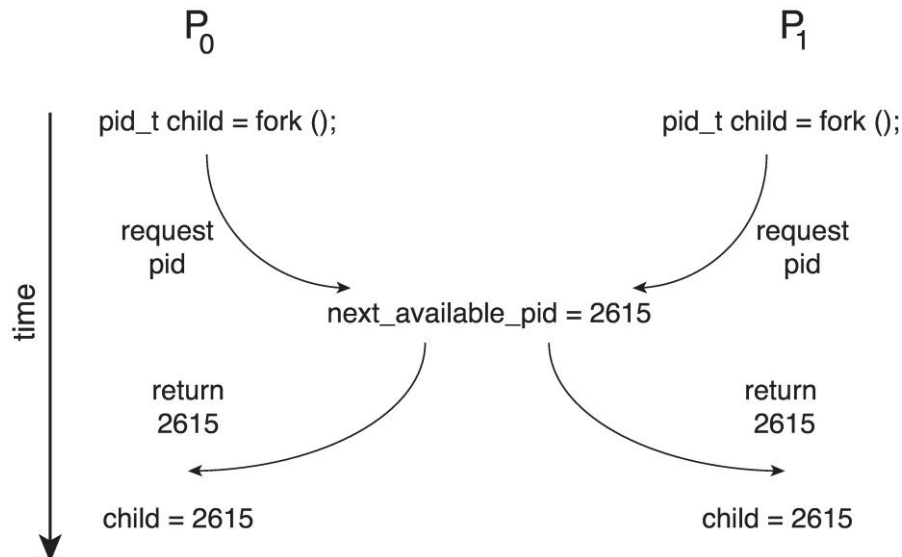
- A situation where several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place is called a **race condition**





# Race Condition in Kernel Code

- Processes  $P_0$  and  $P_1$  are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (pid)



- The same pid could be assigned to two different processes!





# How to avoid race condition?

---

- ❑ Disable interrupts when shared data is modified?
- ❑ Preemptive vs non-preemptive kernel
- ❑ User mode vs Kernel mode
- ❑ Single processor vs Multiprocessors





# Critical-Section Problem

- Consider a system of  $n$  processes  $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has a segment of code, called a **critical section (CS)**, in which the process may be accessing and updating shared data
  - When one process is executing in its CS, no other process is allowed to execute in its CS
- **Critical-Section Problem** is to design a protocol to solve this
  - Each process must ask permission to enter critical section in **entry section**
  - The critical section may be followed by an **exit section**, then **remainder section**

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

General structure of a typical process







# Solution to Critical-Section Problem

---

A solution to the critical-section problem must satisfy the following three requirements:

1. **Mutual Exclusion** - If process  $P_i$  is executing in its CS, then no other processes can be executing in their CSs
2. **Progress** - If no process is executing in its CS and one or more processes wish to enter their CSs, then one of these processes must be allowed to enter its CS
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their CSs after a process has made a request to enter its CS and before that request is granted





# Peterson's Solution

- A software-based solution to the critical-section problem
- A system of two processes  $P_i$  and  $P_j$ , which share two variables:
  - `int turn;`
  - `boolean flag[2]`
- The variable `turn` indicates whose turn it is to enter the critical section
- The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = true` implies that process  $P_i$  is ready.





# Algorithm for Process $P_i$

---

```
while (true){  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /* remainder section */  
  
}
```





# Peterson's Solution (Cont.)

- Provable that the three CS requirements are met:
  1. Mutual exclusion requirement is satisfied

$P_i$  enters CS only if:  
either `flag[j] = false` or `turn = i`
  2. Progress requirement is satisfied
    - ▶ If only one process wants to enter CS, it can enter
    - ▶ If both processes want to enter CS, the variable `turn` determines who can enter next
  3. Bounded waiting requirement is satisfied
    - ▶ A process will enter CS after at most one entry by the other process





# Peterson's Solution (Cont.)

---

- Although useful for demonstrating an algorithm, Peterson's Solution is not guaranteed to work on modern computer architectures
- To improve performance, processors may reorder instructions that have no dependencies
- For single-threaded application this is OK as the result will always be the same
- For multithreaded application the reordering may produce inconsistent or unexpected results!





# Peterson's Solution (Cont.)

---

- Two threads share the data:

```
boolean flag = false;  
int x = 0;
```

- Thread 1 performs

```
while (!flag)  
    ;  
print x
```

- Thread 2 performs

```
x = 100;  
flag = true
```

- What is the expected output?





# Peterson's Solution (Cont.)

- 100 is the expected output.
- However, the operations for Thread 2 may be reordered:

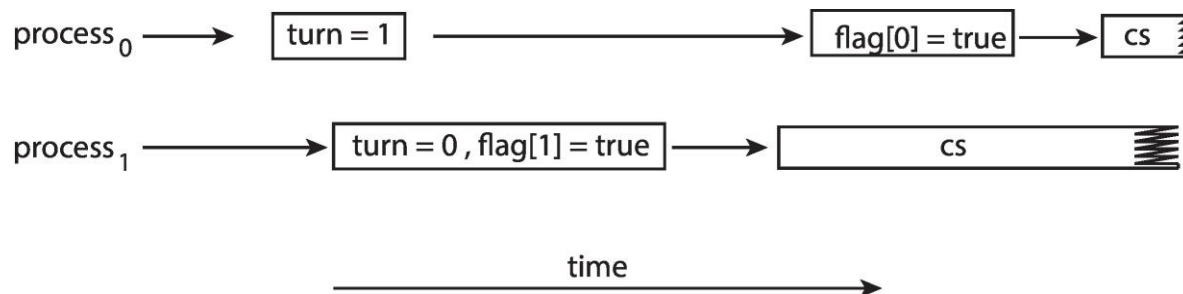
Thread 1 performs

```
while (!flag)
    ;
print x;
```

Thread 2 performs

```
flag = true;
x = 100;
```

- If this occurs, the output may be 0!
- The effects of instruction reordering in Peterson's Solution



- This allows both processes to be in their critical section at the same time!





# Hardware Instructions

---

- Many computer systems provide special **atomic** hardware instructions (**atomic** = non-interruptible) that can be used to solve the critical-section problem
  - **test\_and\_set** instruction: *test-and-modify* the content of a word atomically
  - **compare\_and\_swap** instruction: *swap* the contents of two words atomically







# test\_and\_set Instruction

---

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = true;
    return rv;
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to **true**





# Solution using test\_and\_set()

- Shared boolean variable `lock`, initialized to `false`
- Implementing mutual exclusion:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
    /* remainder section */  
} while (true);
```

Structure of process  $P_i$

The first process that invokes `test_and_set` will set the lock to true, and then enter CS





# compare\_and\_swap Instruction

---

Definition:

```
int compare_and_swap(int *value, int expected, int new_value)
{
    int temp = *value;

    if (*value == expected)
        *value = new_value;
    return temp;
}
```

1. Executed atomically
2. Returns the original value of passed parameter **value**
3. Set the variable **value** to **new\_value** only if **\*value == expected** is true. That is, the swap takes place only under this condition.





# Solution using compare\_and\_swap

- Shared integer variable `lock` initialized to 0;
- Implementing mutual exclusion:

```
while (true){  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
}
```

Structure of process  $P_i$

The first process that invokes `compare_and_swap` will set the lock to 1, and then enter CS





# Mutex Locks

- ❑ Hardware-based solutions to critical section problem are complicated and generally inaccessible to application programmers
- ❑ OS designers build higher-level software tools to solve critical section problem
- ❑ Simplest tool is **mutex lock**
  - ❑ Protect a critical section by first **acquire()** a lock and then **release()** the lock

```
while (true) {  
    acquire lock  
  
    critical section  
  
    release lock  
  
    remainder section  
}
```





# Mutex Lock Definitions

- A mutex lock has a boolean variable **available** indicating if lock is available or not
  - **acquire()** succeeds if lock is available
  - If lock is not available, **acquire()** blocks until lock is released

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}
```

```
release() {  
    available = true;  
}
```

- Calls to **acquire()** and **release()** must be atomic
  - Usually implemented via hardware atomic instructions such as compare-and-swap
- This solution requires **busy waiting**; this lock therefore called a **spinlock**





# Spinlocks

---

- Spinlocks waste CPU cycles
  - Not appropriate for single-processor systems
  - Can avoid busy waiting by putting the waiting process to sleep and then awakening it once the lock becomes available
- Spinlocks have the advantage of not requiring context switch
  - On a multicore system, one thread can spin on one processing core while another thread performs its critical section on another core
  - On multicore systems, spinlocks are the preferable choice for locking if a lock is to be held for a short duration (i.e., less than 2 context switches)





# Semaphore

- A synchronization tool that provides more sophisticated ways (than mutex locks) for process to synchronize their activities
- A semaphore **S** is an integer variable that can only be accessed via two atomic (i.e., non-interruptible) operations

- **wait()** and **signal()**

- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```







# Semaphore Usage

- Two types of semaphores
  - **Counting semaphore** – integer value can range over an unrestricted domain
  - **Binary semaphore** – integer value can range only between 0 and 1
    - ▶ Same as a **mutex lock**
- Semaphores can be used to solve various synchronization problems
- Example 1: a counting semaphore S can be used to control access to a given resource consisting of N instances
  - S is initialized to N

**Structure of a process:**

**wait(S) ;**

**use resource**

**signal(S) ;**





# Semaphore Usage

- Example 2: Consider  $P_1$  and  $P_2$  that require statement  $S_1$  to happen before statement  $S_2$

Create a semaphore “**synch**” initialized to 0

**P1:**

$S_1;$

**signal (synch) ;**

**P2:**

**wait (synch) ;**

$S_2;$





# Semaphore Usage

---

- Example 3: use semaphore for mutual exclusion

Create a semaphore `sem_CS` initialized to 1.

**Structure of a process:**

```
do {  
    wait(sem_CS);  
    Critical section  
    signal(sem_CS);  
    Remainder section  
} while (true);
```





# Semaphore Implementation

- The wait operation suffers from busy waiting

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- To eliminate busy waiting, we associate each semaphore with a waiting queue

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

- When a process executes `wait()` and finds that the semaphore value is not positive, the process is added to the waiting queue
- A `signal()` operation removes one process from the waiting queue and awakens that process





# Semaphore Implementation

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        sleep(); //suspends the process  
    }  
}
```

Can use a FIFO queue to implement the list of waiting processes to ensure bounded waiting

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P); //resumes the operation of P  
    }  
}
```

If semaphore value is negative, its magnitude = number of processes waiting on the semaphore





# Semaphore Implementation

---

- ❑ Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time
- ❑ The implementation becomes the critical section problem where the **wait** and **signal** code are placed in critical section
  - ❑ Can use **test\_and\_set** or **compare\_and\_swap** to implement critical section
  - ❑ OK to have **busy waiting** in critical section implementation because critical section is short






# Problems with Semaphores

- ❑ Incorrect use of semaphore operations:
  - ❑ `signal (mutex)` Critical Section `wait (mutex)`  
Mutual exclusion requirement is violated!
  - ❑ `wait (mutex)` Critical Section `wait (mutex)`  
Process will permanently block on the second call to `wait()`
  - ❑ Omitting `wait (mutex)` or `signal (mutex)`  
Either mutual exclusion is violated or other processes can never enter critical section
- ❑ These – and others – are examples of what can occur when semaphores are used incorrectly





# Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- A **monitor type** is an abstract data type that encapsulates data with a set of functions to operate on that data
  - The set of operations are provided with mutual exclusion; only one process may be active within the monitor at a time
  - Provided in C#, Java
- Pseudocode syntax of a monitor:

```
monitor monitor-name
{
    // shared variable declarations
    function P1 (...) { ... }

    function P2 (...) { ... }

    function Pn (...) {.....}


    initialization code (...) { ... }
}
```







# Schematic view of a Monitor



At any moment, at most one process  
can be active in the monitor





# A Monitor Example

**monitor** SharedAccount

```
{  
    int balance;  
  
    void credit (int amount) {  
        balance = balance + amount  
    }  
  
    void debit (int amount) {  
        balance = balance - amount;  
    }  
  
    initialization_code() {  
        balance = 100;  
    }  
}
```

**SharedAccount.credit(100)** **P1**

**SharedAccount.debit(50)** **P2**



Note: Only one process in monitor at a time



# Condition Variables

- We can use condition variables in monitors to implement more sophisticated synchronization

- Define a condition variable:


**condition *x*, *y*;**

- Two operations are allowed on a condition variable:
  - ***x.wait()*** – a process that invokes the operation is suspended
  - ***x.signal()*** – resumes one of the processes (if any) that invoked ***x.wait()***
    - ▶ If no process is suspended on ***x***, then it has no effect
    - ▶ In case of semaphore, the state is changed on **signal()**





# Monitor with Condition Variables



- There could be processes waiting to enter a monitor
- In addition, there could be processes blocked in queues associated with condition variables





# Condition Variable Example

**monitor** SharedAccount

```
{  
    int balance;  
    condition c;  
  
    void credit (int amount) {  
        balance = balance + amount;  
        c.signal();  
    }  
  
    void debit (int amount) {  
        while (balance - amount < 0)  
            c.wait();  
        balance = balance - amount;  
    }  
  
    initialization_code() {  
        balance = 100;  
    }  
}
```

**SharedAccount.debit(250)** P1

**SharedAccount.credit(100)** P2

**SharedAccount.credit(50)** P3





# Condition Variables Choices

- If process P invokes **`x.signal()`** , and process Q is suspended on **`x`**, what should happen next?
  - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options:
  - **Signal and continue** – Q waits until P either leaves the monitor or waits for a condition
  - **Signal and wait** – P waits until Q either leaves the monitor or waits for a condition
  - A compromise: P executing **`signal()`** immediately leaves the monitor and Q is resumed





# Deadlocks

- ❑ **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused only by one of the waiting processes
- ❑ Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$	$P_1$
<code>wait(S) ;</code>	<code>wait(Q) ;</code>
<code>wait(Q) ;</code>	<code>wait(S) ;</code>
<code>...</code>	<code>...</code>
<code>signal(S) ;</code>	<code>signal(Q) ;</code>
<code>signal(Q) ;</code>	<code>signal(S) ;</code>

- ❑ Suppose  $P_0$  executes `wait(S)` and  $P_1$  executes `wait(Q)`.
- ❑ When  $P_0$  executes `wait(Q)`, it must wait until  $P_1$  executes `signal(Q)`.
- ❑ Similarly, When  $P_1$  executes `wait(S)`, it must wait until  $P_0$  executes `signal(S)`.
- ❑ Since `signal(S)` and `signal(Q)` can never be executed,  $P_0$  and  $P_1$  are **deadlocked**.

