


Chapter 9: Main Memory


Section 9.1-9.2,9.3.1





Background

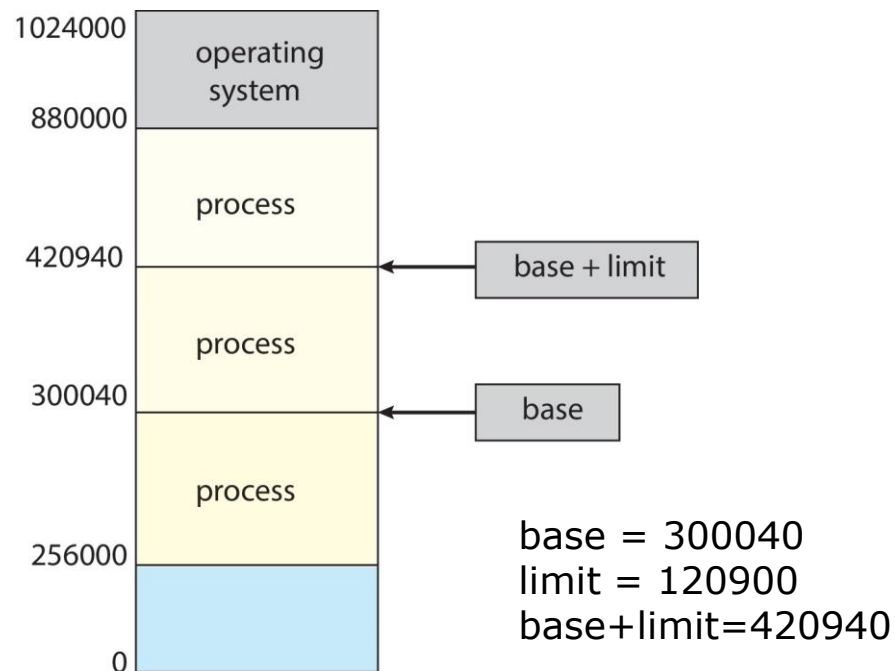
- ❑ Program must be brought (from disk) into memory and placed within a process for it to be run
- ❑ Main memory and registers are only storage CPU can access directly
- ❑ Memory unit only sees a stream of:
 - ❑ read requests & addresses, or
 - ❑ write requests & address + data
- ❑ Register access is done in one CPU cycle (or less)
- ❑ Main memory access can take many CPU cycles, causing processor to **stall**
- ❑ **Cache** is typically added on CPU chip to speedup memory access





Memory Protection


- Need to ensure that a process can access only those memory addresses in its address space
- We can provide this protection using **base** and **limit registers** that define the range of legal memory addresses that a process may access
 - **Base register** holds the smallest legal memory address
 - **Limit register** specifies the size of the range





Memory Protection Hardware

- CPU hardware must compare every memory access generated in user mode with base and limit registers




- Instructions to load base and limit registers are privileged instructions





Address Binding (1)

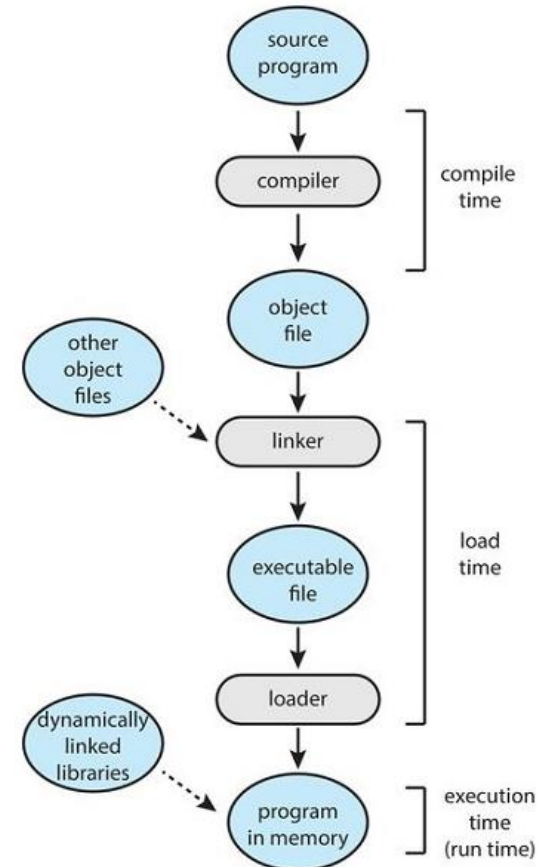
- A user process may reside in any part of the physical memory
 - The first address of a user process need not be 0000
- Addresses represented in different ways at different stages of a program's life
 - Addresses in source program usually symbolic (e.g., int count)
 - A compiler **binds** symbolic addresses to relocatable addresses (e.g. “14 bytes from beginning of this module”)
 - Linker or loader binds relocatable addresses to absolute addresses
 - ▶ e.g., relocatable address 14 is bound to absolute address 74014
 - Each binding maps one address space to another





Address Binding (2)

- ❑ Binding of instructions and data to memory addresses can happen at three different stages:
 - ❑ **Compile time:** If memory location of process known a priori, **absolute code** can be generated;
 - ▶ Must recompile code if starting location changes
 - ❑ **Load time:** Compiler must generate **relocatable code** if memory location of process is not known at compile time; final address binding is delayed until load time
 - ▶ If starting address changes, need reload the user code
 - ❑ **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - ▶ Need hardware support for address mapping (e.g., relocation register)
 - ▶ Used by most operating systems





Logical vs. Physical Address Space


- The concept of a **logical address space** that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes
- Logical and physical addresses differ in execution-time address-binding scheme





Memory-Management Unit (MMU)

- MMU is a hardware device that maps logical addresses to physical addresses at run time




- Different mapping methods possible, covered in the rest of this chapter





Memory-Management Unit (Cont.)


- Consider a simple mapping scheme, which is a generalization of the base-register scheme
- The base register now called **relocation register**
- The value in the relocation register is added to every address generated by a user process at the time it is sent to memory
- The user program deals with *logical* addresses; it never sees the *real* physical addresses
 - Execution-time binding occurs when a reference is made to a location in memory





Contiguous Memory Allocation

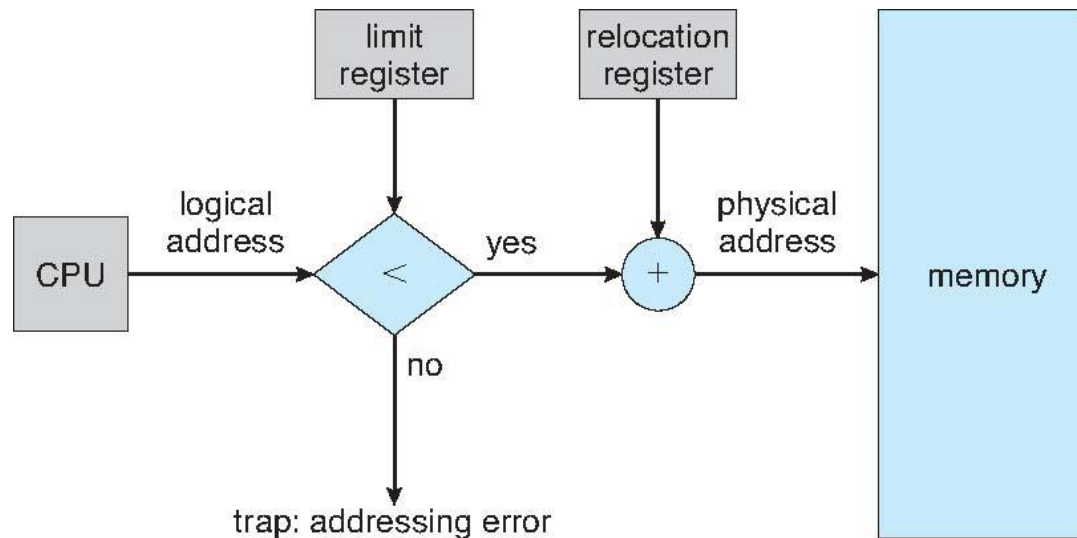
- ❑ Main memory must support both OS and user processes
- ❑ Limited resource, must allocate efficiently
- ❑ Contiguous allocation is one early method
- ❑ Main memory usually divided into two **partitions**: one for OS and one for user processes
 - ❑ OS usually held in high memory
 - ❑ Each user process contained in a single contiguous section of memory





Contiguous Memory Allocation (Cont.)


- Relocation register and limit register used to protect user processes from each other, and from changing OS code and data
 - Relocation register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address to physical address *dynamically*





Memory Allocation

- Initially, all memory is available and is considered one large **hole**
 - A **hole** is a block of available memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - If the hole is too large, the unused part is returned to the set of holes
- When a process terminates, it releases its block of memory
 - If the new hole is adjacent to other holes, these holes are merged to form a larger hole
- In general, holes of various sizes are scattered throughout memory
- Operating system maintains information about a) allocated memory and b) free memory (holes)





Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of holes?

- Commonly used strategies:
 - **First-fit**: Allocate the **first** hole that is big enough
 - **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless ordered by size
 - ▶ Produces the smallest leftover hole
 - **Worst-fit**: Allocate the **largest** hole; must also search entire list, unless ordered by size
 - ▶ Produces the largest leftover hole
- First-fit and best-fit better than worst-fit in terms of storage utilization
- First-fit generally faster than best-fit





Example

Given six memory holes of

300 KB, 600 KB, 350 KB, 200 KB, 750 KB, 125 KB (in order),

how would the first-fit, best-fit, and worst-fit algorithms place a process of size 115 KB?

□ **First fit:** The process is put in 300-KB hole, leaving
185 KB ($=300\text{KB}-115\text{KB}$), 600 KB, 350 KB, 200 KB, 750 KB, 125 KB

□ **Best fit:** 115 KB is put in 125-KB hole, leaving
300 KB, 600 KB, 350 KB, 200 KB, 750 KB, 10 KB

□ **Worst fit:** 115 KB is put in 750-KB hole, leaving
300 KB, 600 KB, 350 KB, 200 KB, 635 KB ($=750\text{KB}-115\text{KB}$), 125 KB


how would the first-fit, best-fit, and worst-fit algorithms place another process of size 200 KB?






Fragmentation


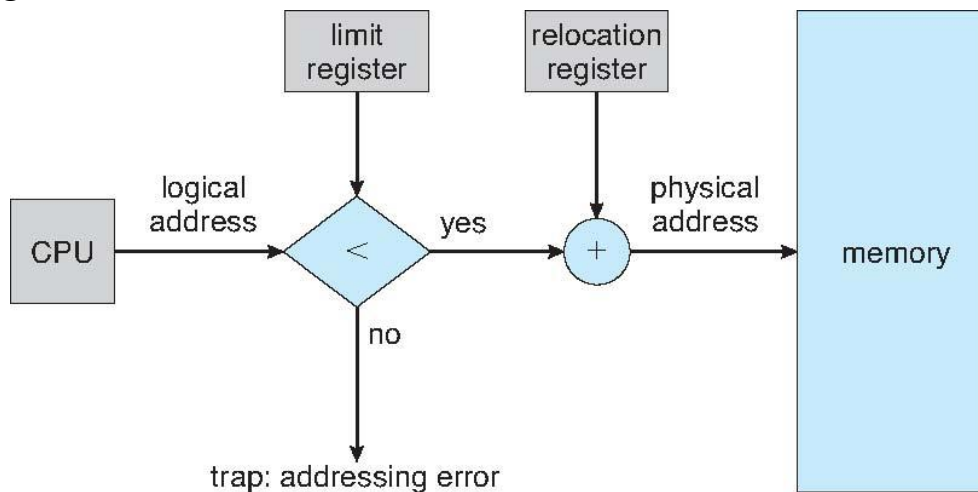
- **External Fragmentation** – there is enough total memory space to satisfy a request, but the available spaces are not contiguous
 - Statistical analysis of first-fit reveals that given N allocated blocks, another $0.5N$ blocks will be lost to external fragmentation
 - That is, $1/3$ of memory may be unusable -> **50-percent rule**
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; the size difference is memory internal to a partition, but not being used
 - E.g., 4000 bytes requested, 4096 bytes (4KB) allocated





Overcoming External Fragmentation

- One solution to external fragmentation is **compaction**
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible *only* if relocation is dynamic, i.e., address binding done at execution time using a relocation register



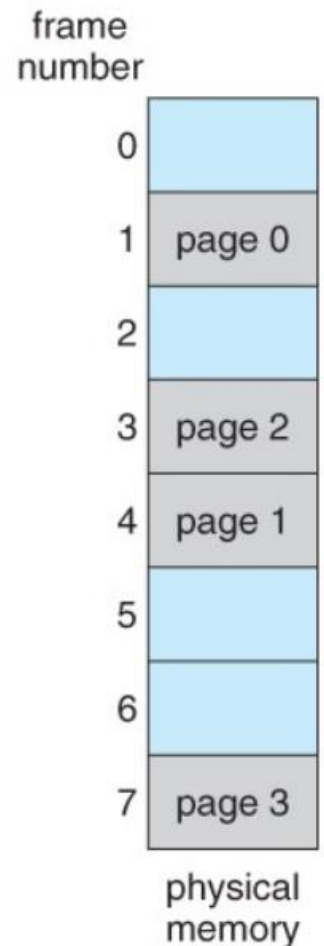
- Another solution is to permit the physical address space of a process to be noncontiguous
 - E.g., paging





Paging


- ❑ Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
- ❑ Divide physical memory into fixed-sized blocks called **frames**
 - ❑ Frame size is power of 2, typically between 4KB (2^{12} bytes) and 1GB (2^{30} bytes)
- ❑ Divide logical memory into blocks of same size called **pages**
- ❑ To run a process of size **N** pages, need to find **N** free frames and load pages of process to frames
- ❑ OS keeps track of all free frames
- ❑ External fragmentation is avoided; still have internal fragmentation
- ❑ Paging is used in most operating systems





Paging


- ❑ Kernel maintains a **page table** for each process
 - ❑ Page table has an entry for each page of the process, entry indicates the memory frame allocated to the page
 - ❑ Page table is used to translate logical addresses to physical addresses





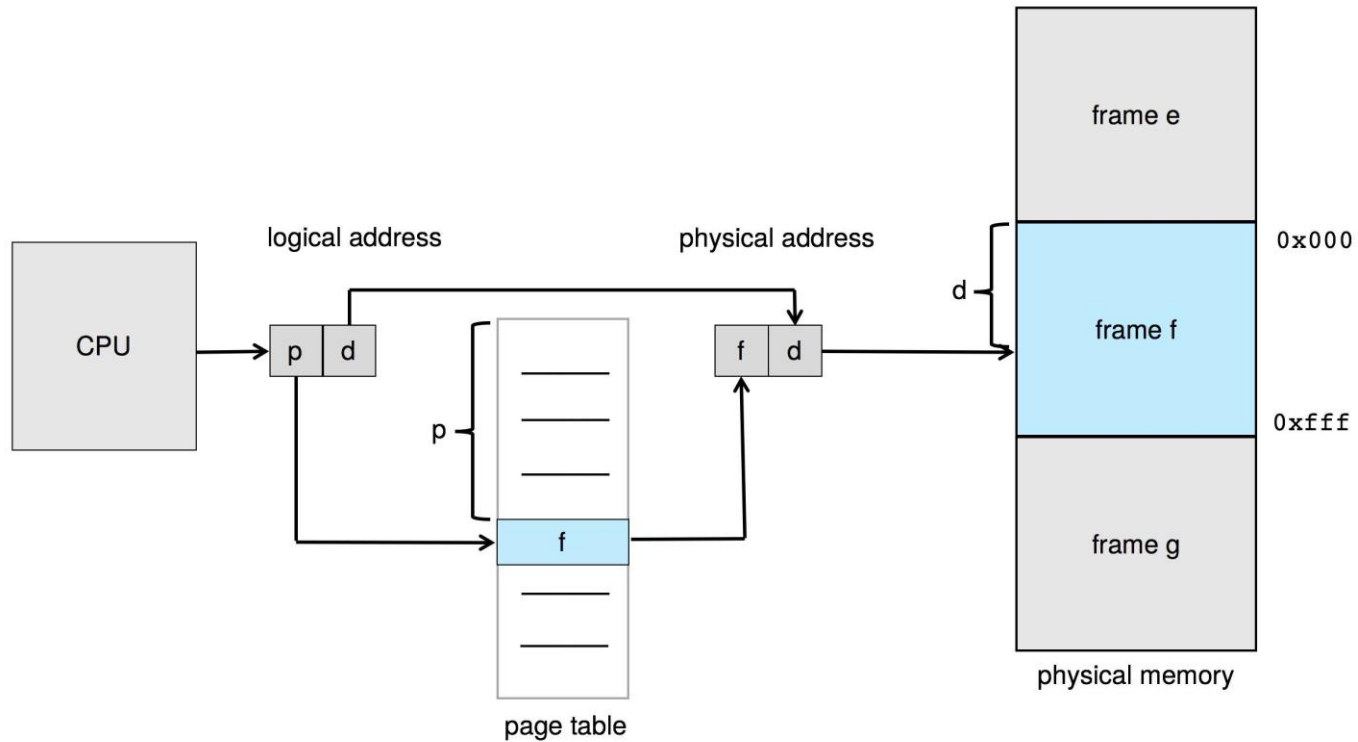
Address Translation Scheme

- Address generated by CPU (i.e., logical address) is divided into:
 - **Page number** (p) – used as an index into a **page table** which contains base address of each frame in physical memory
 - **Page offset** (d) – location in the frame
 - Base address of frame is combined with the page offset to define the physical memory address that is sent to the memory unit
- If size of logical address space = 2^m bytes and page size = 2^n bytes, then
 - High-order $m-n$ bits of logical address represent the page number
 - Low-order n bits represent the page offset





Paging Hardware



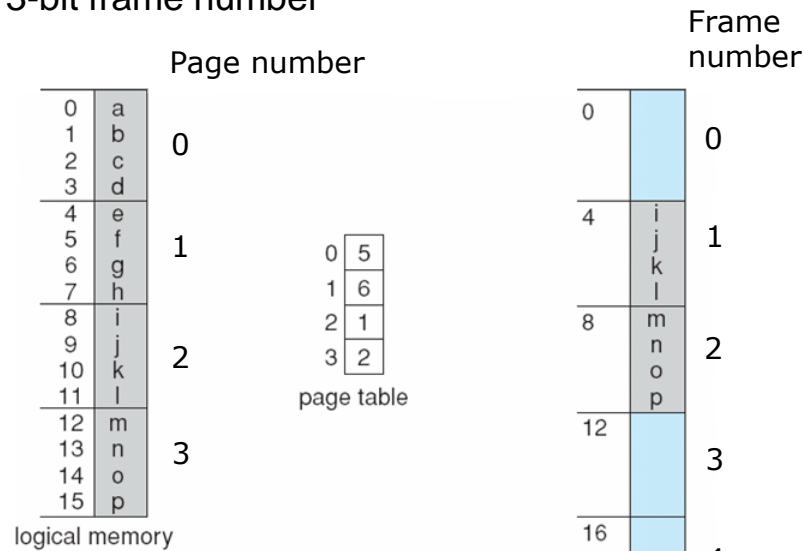
MMU translates a logical address to a physical address





Paging Example

- Page size = 4 bytes, physical memory = 32 bytes (i.e., 8 frames), logical memory = 16 bytes (i.e., 4 pages)
 - 2-bit page offset ($n = 2$)
 - 4-bit logical address ($m = 4$), 2-bit page number ($m - n = 2$)
 - 5-bit physical address, 3-bit frame number



- Logical address 3 (page 0, offset 3) maps to physical address $5 \times 4 + 3 = 23$
 - In binary, logical address **00**11 maps to physical address **10**111
- Logical address 13 (page 3, offset 1) maps to physical address $2 \times 4 + 1 = 9$
 - In binary, logical address **11**01 maps to physical address **01**001





Internal Fragmentation in Paging


- Calculating internal fragmentation
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
- Worst case happens when a process needs n pages plus 1 byte → internal fragmentation is almost 1 frame
 - On average internal fragmentation = $1/2$ frame size
- So small page sizes desirable?
 - But each page table entry takes memory to track
 - Page sizes have grown over time
 - Today page sizes typically 4KB or 8KB






Free Frames

- OS keeps track of free frames



(a)

Before allocation



(b)

After allocation

