


Chapter 14: File System Implementation

Section 14.1,14.4-14.5.1,14.5.2





File-System Structure

- **File system** resides on secondary storage (disks)
 - Provides efficient and convenient access to disk by allowing data to be stored, located, and retrieved easily
- Disk provides in-place rewrite and random access
 - I/O transfers performed in units of **blocks**
 - Each block has one or more **sectors** (usually 512 bytes)
- File system translates file name into file number and location by maintaining directory structure and **file control block (FCB)** (**inodes** in UNIX)





File Control Block

- Per-file **File Control Block (FCB)** contains many details about the file
 - typically inode number, permissions, size, dates

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks





In-Memory File System Structures

- **system-wide open-file table** contains a copy of the FCB of each file and open count
- **per-process open-file table** contains pointers to appropriate entries in system-wide open-file table as well as other info (e.g., file pointer, access mode)

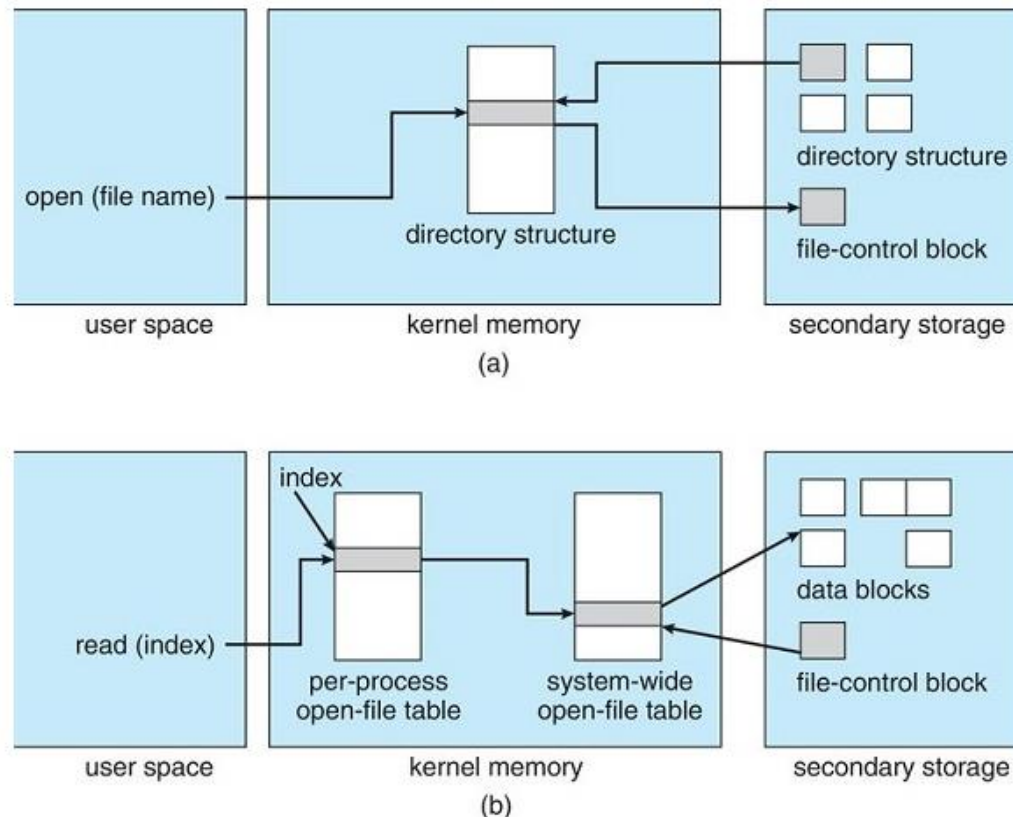


Figure 14.3 In-memory file-system structures. (a) File open. (b) File read.





File Structure

- A file is composed of contiguous **logical blocks**
 - Each logical block has a fixed size
- Secondary storage is composed of **physical blocks**, each having the same size as a logical block
- A file's logical blocks are mapped to physical blocks





Allocation Methods


- An allocation method determines how disk blocks are allocated to files
 - Should utilize disk space efficiently and allow files to be accessed quickly
- Three allocation methods
 - **Contiguous allocation**
 - **Linked allocation**
 - **Indexed allocation**





Contiguous Allocation

- Each file occupies a set of contiguous disk blocks
 - First-fit or best-fit
 - Simple – only starting location (block #) and length (number of blocks) are required
 - Efficient for both sequential access and direct access
- Problems
 - External fragmentation, need for compaction
 - File cannot grow, need to know file size



directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2





Extent-Based Systems


- Many newer file systems (i.e., Veritas File System) use a modified contiguous allocation scheme
- An **extent** is a contiguous chunk of space
- An extent is allocated to a file initially
 - If the amount is not large enough, another extent is added
- A file consists of one or more extents
 - For each extent, the location of the first block and the block count are recorded





Linked Allocation

- ❑ **Linked allocation** – each file is a linked list of blocks, blocks may be scattered anywhere on the disk
 - ❑ Each block contains pointer to next block
 - ❑ File ends at nil pointer
- ❑ No external fragmentation
- ❑ A file can continue to grow as long as free blocks are available
- ❑ Efficient for sequential access
- ❑ Inefficient for direct access – access to i^{th} block requires i disk reads
- ❑ Overhead - space required for pointers
- ❑ Reliability problem - what if a pointer is damaged?







File-Allocation Table (FAT)

- ❑ **FAT** is an variation of linked allocation
- ❑ Beginning of volume has a file-allocation table having one entry for each disk block
 - ❑ Entry for block x contains the block number of the next block in file
 - ❑ Table value 0 indicates unused block
- ❑ New block allocation is simple
- ❑ Efficient direct access if FAT is loaded into the memory

directory entry






Indexed Allocation

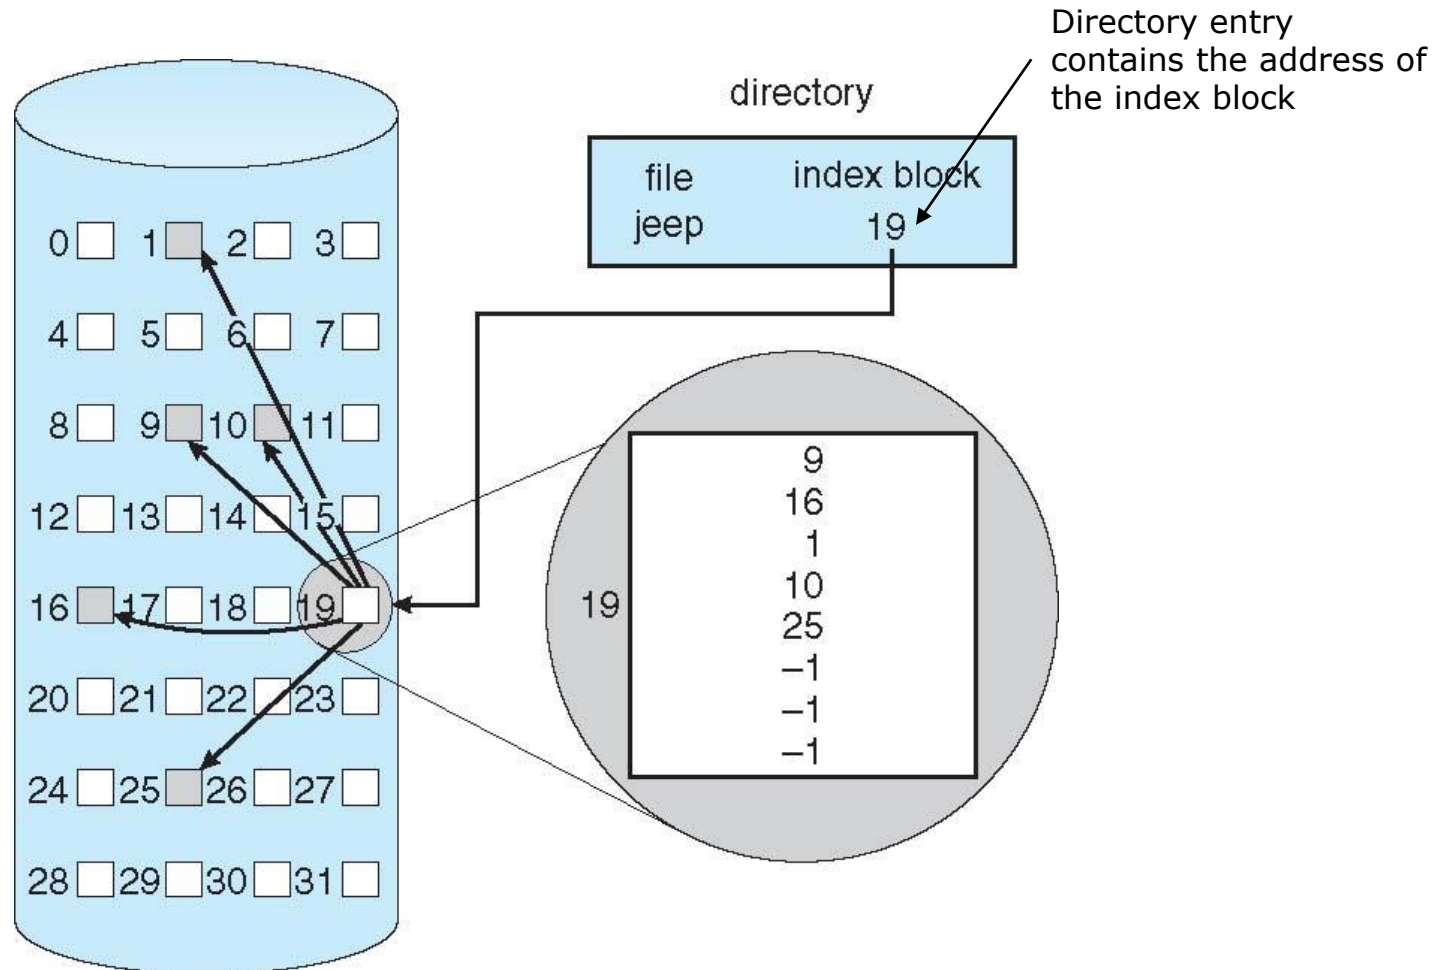
□ Indexed allocation

- Each file has its own **index block** of pointers to its data blocks
- To find the i^{th} block of the file, we use the pointer in the i^{th} index block entry





Example of Indexed Allocation





Indexed Allocation (Cont.)

- ❑ Efficient direct access if index block is in memory
- ❑ File size can change
- ❑ No external fragmentation, but have overhead of index block


- ❑ If block size is 1KB and pointer size is 4 bytes, how many data blocks can a file have?
 - ❑ 256 data blocks → max file size = 256KB
- ❑ What if the size of a file exceeds 256KB?
 - ❑ Solution 1: Linked index
 - ❑ Solution 2: Multilevel index
 - ❑ Solution 3: Combined scheme





Linked Index Allocation


- To allow large files, we can link several index blocks together
- An index block contains
 - A small header giving name of the file
 - Set of N block addresses
 - Pointer to next index block





Multilevel Index Allocation

- Two-level index: A first-level index block points to a set of second-level index blocks, which in turn point to file blocks
- If block size is 4KB and pointers are 4 bytes, what is the maximum allowed file size?
 - $1,024 * 1024$ data blocks and file size of up to 4GB

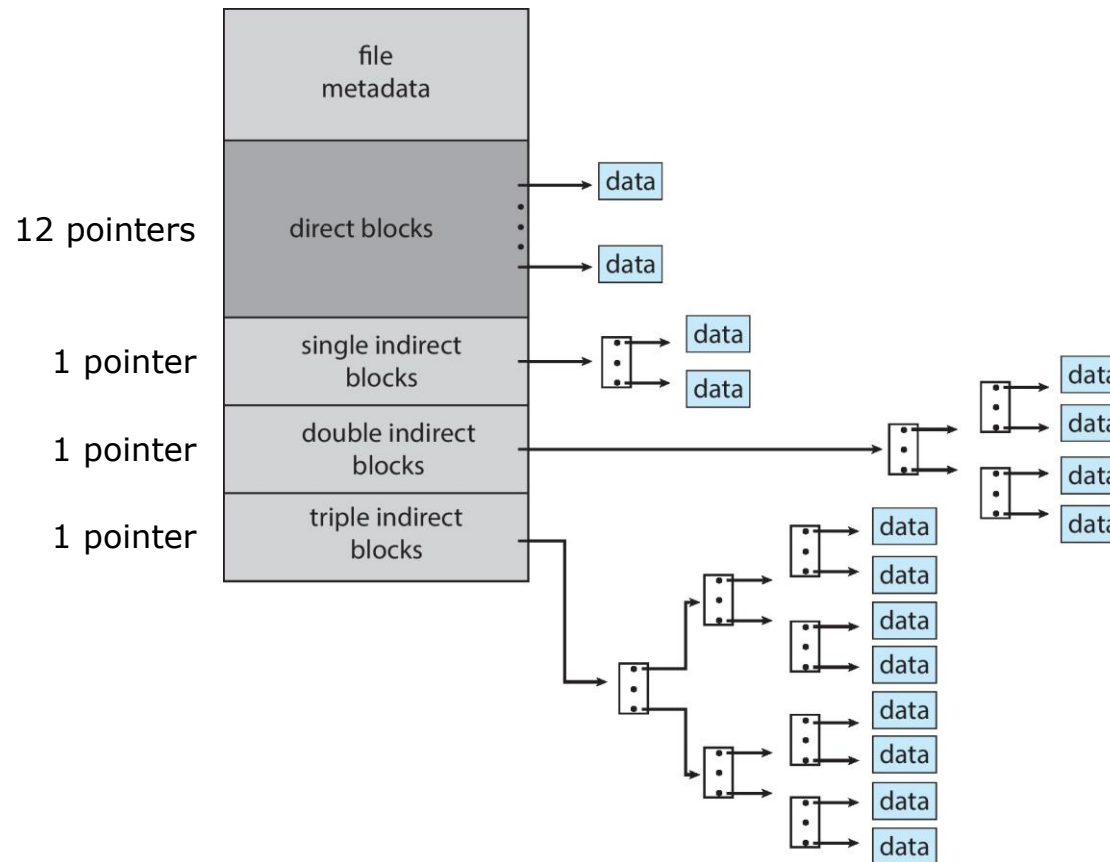




Combined Scheme: UNIX

4KB block size, 32-bit pointers

Number of blocks that can be allocated to a file = $12 + 1024 + 1024^2 + 1024^3$



The UNIX inode





Performance

- ❑ Best method depends on file access type
- ❑ Contiguous allocation great for sequential and direct accesses
 - ❑ Only one access is required to get a block
- ❑ Linked allocation good for sequential access, not direct access
 - ❑ Access to the *i*th block requires *i* block reads
- ❑ Some systems support direct-access files by using contiguous allocation and sequential-access files by using linked allocation
 - ❑ Declare access type at file creation
- ❑ Indexed allocation more complex
 - ❑ Single-level index: access requires index block read then data block read
 - ❑ Two-level index: access requires 2 index block reads then data block read
- ❑ Some systems combine contiguous allocation with indexed allocation
 - ❑ Using contiguous allocation for small files (up to 4 blocks) and switching to indexed allocation if the file grows large





Free-Space Management

- File system maintains **free-space list** to track available blocks
- **Bit vector** or **bitmap** (n blocks)

0	1	2	3				$n-1$
0	0	1	1	1	0	...	1

$$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{block}[i] \text{ free} \\ 0 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

The 1st non-0 word is scanned for the first 1 bit, which is the location of the first free block

Block number calculation:

(number of bits per word) * (number of 0-value words) +
offset of first 1 bit

CPUs have instructions to return offset within word of first “1” bit





Free-Space Management (Cont.)

- Bitmap requires extra space

- Example:

block size = 4KB = 2^{12} bytes

disk size = 2^{40} bytes (1 terabyte)

number of blocks $n = 2^{40}/2^{12} = 2^{28} \rightarrow$ need a bit map
of 2^{28} bits (or 32MB) to track the free blocks

if clusters of 4 blocks \rightarrow 8MB of memory

- Can get n consecutive free blocks easily

0	1	2	3				$n-1$
0	0	1	1	1	0	...	1





Linked Free-Space List on Disk

- Link together all the free blocks
 - Allocating a free block to a file is easy
 - Cannot get contiguous space easily

free-space list head

