


# Chapter 17: Protection

Section 17.1-17.2,17.4-17.6





# Goals of Protection

---

- A computer system consists of a collection of objects
  - **Hardware objects** (such as devices) and **software objects** (such as files, programs, semaphores)
- Each object has a unique name and can be accessed through a well-defined set of operations
- Protection involves controlling the access of processes and users to the resources defined by a computer system
  - OS ensures that each object is accessed correctly and only by those processes that are allowed to do so
- Protection mechanisms provide a means for specifying the controls to be imposed and a means of enforcement





# Principles of Protection

---

- Guiding principle – **principle of least privilege**
  - Processes and users should be given just enough **privileges** to perform their tasks
- **Compartmentalization** – a derivative of the principle of least privilege
  - Process of protecting each individual system component through the use of specific permissions and access restrictions
  - Properly set **permissions** can limit damage if entity has a bug, gets abused





# Domain of Protection

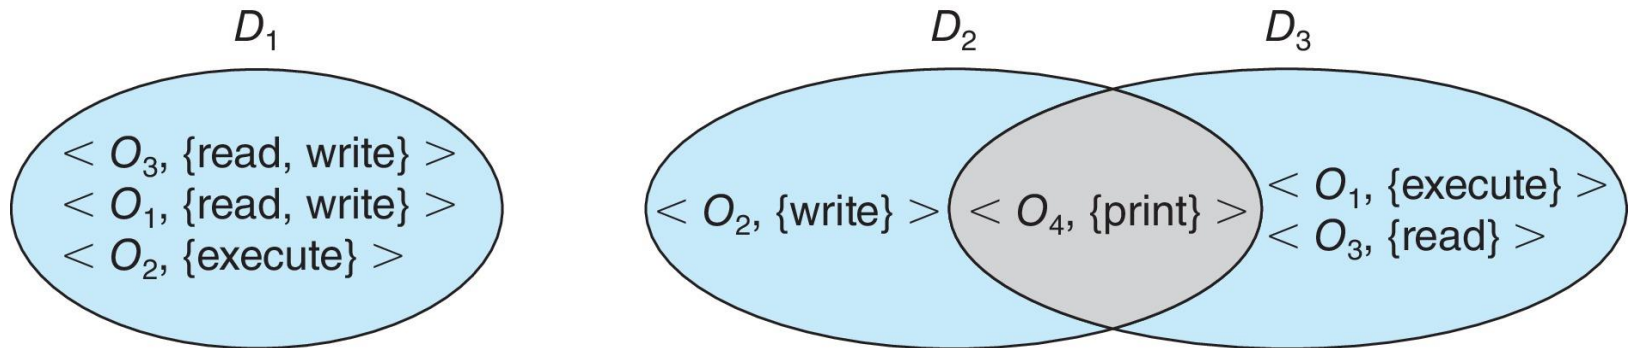
- ❑ Process should only have access to objects it currently requires to complete its task – the **need-to-know** principle
- ❑ Implementation can be via process operating in a **protection domain**
  - ❑ Specifies resources process may access
  - ❑ Each domain specifies a set of objects and types of operations that may be invoked on each object
  - ❑ Ability to execute an operation on an object is an **access right**
    - ▶  $\text{access-right} = \langle \text{object-name}, \text{rights-set} \rangle$ , where *rights-set* is a subset of all valid operations that can be performed on the object
  - ❑ A domain is a set of access-rights





# Domain Structure

- Domain = set of access-rights
- Domains may share access rights
- Association between a process and a domain can be **static** or **dynamic**
  - If dynamic, a process can switch from one domain to another





# Domain Structure

---

A domain can be realized in several ways:

- Each user may be a domain
  - The set of objects that can be accessed depends on the identify of the user
- Each process may be a domain
  - The set of objects that can be accessed depends on the identify of the process
- Each procedure may be a domain
  - The set of objects that can be accessed corresponds to the local variables defined within the procedure





# Access Matrix

- **Access matrix** provides a mechanism for protection
  - Rows represent domains
  - Columns represent objects
  - Entry `access(i, j)` defines the set of operations that a process executing in domain  $D_i$  can invoke on object  $O_j$

object domain	$F_1$	$F_2$	$F_3$	printer
$D_1$	read		read	
$D_2$				print
$D_3$		read	execute	
$D_4$	read write		read write	






# Use of Access Matrix

- If a process in domain  $D_i$  tries to do “op” on object  $O_j$ , then “op” must be in the entry **access**( $i, j$ )
- User who creates an object can define access column for that object
- We can control domain switching by including domains among the objects of the access matrix
  - Switching from domain  $D_i$  to domain  $D_j$  is allowed if and only if the access right **switch** is in entry **access**( $i, j$ )

domain \ object	$F_1$	$F_2$	$F_3$	laser printer	$D_1$	$D_2$	$D_3$	$D_4$
$D_1$	read		read			switch		
$D_2$				print			switch	switch
$D_3$		read	execute					
$D_4$	read write		read write		switch			







# Use of Access Matrix

---

- Processes can change contents of access-matrix entries
- Special operations: copy, owner, control
  - **Copy**: ability to copy an access right from one domain to another
  - **Owner**: ability to add or remove access rights
  - **Control**: ability to change the entries in a row
  - **Copy** and **Owner** rights applicable to an object
  - **Control** right applicable to a domain





# Access Matrix with Copy Rights

- Ability to *copy* an access right from one domain to another is denoted by an asterisk(\*) appended to the access right
  - Access right can only be copied within column

object \ domain	$F_1$	$F_2$	$F_3$
$D_1$	execute		write*
$D_2$	execute	read*	execute
$D_3$	execute		

(a)

A process executing in domain  $D_2$  can modify the access matrix of (a) to the access matrix of (b)

object \ domain	$F_1$	$F_2$	$F_3$
$D_1$	execute		write*
$D_2$	execute	read*	execute
$D_3$	execute	read	

(b)





# Access Matrix With Owner Rights

- If **access**(*i*, *j*) includes the *owner* right, then a process executing in  $D_i$  can add and remove any right in any entry in column *j*

object \ domain	$F_1$	$F_2$	$F_3$
$D_1$	owner execute		write
$D_2$		read* owner	read* owner write
$D_3$	execute		

(a)

The access matrix of (a) can be modified to the access matrix shown in (b)

object \ domain	$F_1$	$F_2$	$F_3$
$D_1$	owner execute		write
$D_2$		owner read* write*	read* owner write
$D_3$		write	write

(b)





# Access Matrix With *Control* Rights

- If **access** ( $i, j$ ) includes the *control* right, then a process executing in  $D_i$  can modify any right from row  $D_j$

object \ domain	$F_1$	$F_2$	$F_3$	laser printer	$D_1$	$D_2$	$D_3$	$D_4$
$D_1$	read		read			switch		
$D_2$				print			switch	switch
$D_3$		read	execute					
$D_4$	read write		read write		switch			

If we include control right here, then a process executing in  $D_2$  can modify  $D_4$ , as shown in the lower figure

object \ domain	$F_1$	$F_2$	$F_3$	laser printer	$D_1$	$D_2$	$D_3$	$D_4$
$D_1$	read		read			switch		
$D_2$				print			switch	switch control
$D_3$		read	execute					
$D_4$	write		write		switch			





# Use of Access Matrix (Cont.)

---

- **Access matrix** design separates mechanism from policy
  - Mechanism
    - ▶ Operating system provides access matrix + rules
    - ▶ It ensures that the matrix is only manipulated by authorized agents and that rules are strictly enforced
  - Policy
    - ▶ System designers and users dictate policy
    - ▶ Which domains can access which objects in which ways





# Implementation of Access Matrix

- Generally, a sparse matrix
- Option 1 – Global table
  - Store ordered triples **<domain, object, rights-set>** in table
    - ▶ E.g.,  $\langle D_1, F_1, \{\text{read}\} \rangle$ ,  $\langle D_4, F_3, \{\text{read}, \text{write}\} \rangle$
  - A requested operation  $op$  on object  $O_j$  within domain  $D_i \rightarrow$  search table for triple  $\langle D_i, O_j, R \rangle$  with  $op \in R$ 
    - ▶ If triple is found, the operation is allowed
    - ▶ Otherwise, an exception condition is raised
  - But table could be large  $\rightarrow$  won't fit in main memory
  - Difficult to group objects or domains (consider an object that all domains can read)

domain \ object	$F_1$	$F_2$	$F_3$	laser printer	$D_1$	$D_2$	$D_3$	$D_4$
$D_1$	read		read			switch		
$D_2$				print			switch	switch
$D_3$		read	execute					
$D_4$	read write		read write		switch			





# Implementation of Access Matrix (Cont.)

- Option 2 – Access lists for objects
  - Each column implemented as an access list for one object
  - Per-object list consists of ordered pairs **<domain, rights-set>** defining all domains with non-empty set of access rights for the object
    - ▶ E.g., access list for  $F_1 = [<D_1, \{read\}>, <D_4, \{read, write\}>]$
  - Operation  $op$  on object  $O_j$  is attempted in domain  $D_i \rightarrow$  search the access list for  $O_j$  for entry  $<D_i, R>$  with  $op \in R$ 
    - ▶ If entry is found, the operation is allowed. Otherwise, access is denied
  - Can be extended to contain default set of access rights  $\rightarrow$  If  $op \in$  default set, also allow access

object domain \	$F_1$	$F_2$	$F_3$	laser printer	$D_1$	$D_2$	$D_3$	$D_4$
$D_1$	read		read			switch		
$D_2$				print			switch	switch
$D_3$		read	execute					
$D_4$	read write		read write		switch			





# Implementation of Access Matrix (Cont.)

- ❑ Option 3 – Capability list for domains
  - ❑ Instead of object based, list is domain based
  - ❑ **Capability list** for a domain is a list of objects together with operations allowed on those objects
    - Object F1 – read
    - Object F4 – read, write, execute
    - Object F5 – read, write, delete, copy
  - ❑ Object represented by its name or address, called a **capability**
  - ❑ To execute operation  $op$  on object  $O_j$ , process requests operation  $op$  and specifies capability (or pointer to capability) for  $O_j$  as a parameter
    - ▶ Possession of capability means access is allowed
  - ❑ Capability list associated with domain but never directly accessible to a process executing in that domain
    - ▶ Capability list is a protected object, maintained by OS and accessed by user indirectly
    - ▶ Capabilities like a “secure pointer”







# Capability-Based Systems

- ❑ Capabilities are used in Linux and Android
  - ❑ Essentially slices up root privileges into distinct areas, each represented by a bitmap bit
  - ❑ Fine grain control over privileged operations can be achieved by setting or masking the bitmap
    - ▶ Process or thread starts with the full set of permitted capabilities, voluntarily decreases set during execution
    - ▶ Essentially a direct implementation of the principle of least privilege
- ❑ List of Linux capabilities can be found at <http://man7.org/linux/man-pages/man7/capabilities.7.html>



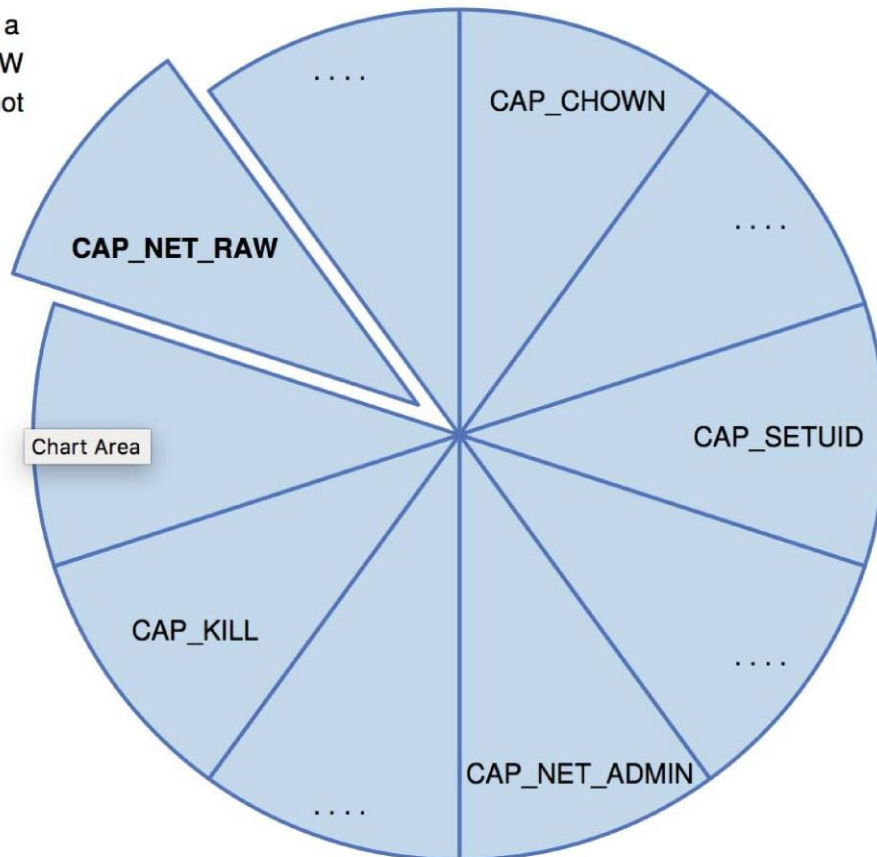


# Capabilities in Linux

In the old model, even a simple ping utility would have required root privileges, because it opens a raw (ICMP) network socket

Capabilities can be thought of as "slicing up the powers of root" so that individual applications can "cut and choose" only those privileges they actually require

With capabilities, ping can run as a normal user, with CAP\_NET\_RAW set, allowing it to use ICMP but not other extra privileges





# Comparison of Implementations

- Many trade-offs to consider
  - Global table is simple, but can be large
  - Access lists correspond to needs of users
    - ▶ When a user creates an object, he can specify which domains can access the object, and what operations are allowed
    - ▶ Determining set of access rights for each domain is difficult as information is non-localized
    - ▶ Every access to an object must be checked
      - When access list is long, search will be slow
  - Capability lists useful for localizing information for a given process
    - ▶ Process attempting access must present a capability for that access





# Comparison of Implementations (Cont.)

- Most systems use combination of access lists and capabilities
  - First access to an object → access list searched
    - ▶ If allowed, capability created and attached to process
      - Additional accesses need not be checked
    - ▶ After last access, capability destroyed
  - Consider file system with an access list per file
    - ▶ When a process opens a file, access permission is checked
    - ▶ If access is allowed, per-process open-file table entry created, entry has capability for allowed operations (e.g., read-only, write-only, read-write)

